

Spark streaming

Igor Yakushin
`ivy2@uchicago.edu`

December 27, 2019

Streaming

- Streaming allows to process continuously arriving data
- There are two streaming interfaces in Spark:
 - **Structured Streaming** - DataFrame API
 - **Spark Streaming (DStreams)** - RDD API
- We shall only consider Structured Streaming.
- Structured Streaming is a scalable and fault-tolerant stream processing engine built on the Spark SQL engine.
- One can express streaming computation the same way one would express a batch computation on static data.
- The Spark SQL engine will take care of running it incrementally and continuously and updating the final result as streaming data continues to arrive
- Structured Streaming can use as sources: files, Kafka, socket (for testing)

Streaming: Lab 6: word_count.py

```
from pyspark.sql.functions import explode, split
import sys
port = int(sys.argv[1]); host = sys.argv[2]
spark = SparkSession.builder.getOrCreate()
lines = spark.readStream.format("socket") \
    .option("host",host).option("port",port).load()
words = lines.select(
    explode(
        split(lines.value, " ")
    ).alias("word")
)
wordCounts = words.groupBy("word").count()
query = wordCounts \
    .writeStream \
    .outputMode("complete") \
    .format("console") \
    .start()
query.awaitTermination()
```

Streaming: Lab 6: word_count.py

- Generate random port number from 9000 to 10000 and write the corresponding commands used below into the files `mync.sh` and `mystream.sh`.
- Open the second terminal on the same host and run `netcat` in it:
`source env.sh`
`make nc`
- `netcat` will connect to `localhost:<port>`
- In the first terminal, start pyspark streaming program:
`make word_count`
- Now pyspark listens to input on `localhost:<port>`
- In the second terminal type some words. Once you hit `Enter`, all the words including the current ones are reprocessed by pyspark. Keep entering new lines and observe the output in the first terminal.
- To stop both processes, type `Ctrl+C` in the window with pyspark.

Streaming: Lab 6: age.py

```
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType
from pyspark.sql.functions import avg

spark = SparkSession.builder.getOrCreate()

userSchema = StructType().add("name", "string")\
                        .add("age", "integer")
csvDF = spark.readStream.option("sep", ";") \
    .schema(userSchema).csv("input_csv")

averageAge = csvDF.select(avg(csvDF.age))

query = averageAge.writeStream.outputMode("complete")\
    .format("console").start()
query.awaitTermination()
```

Streaming: Lab 6: age.py

- In this lab data is processed as files are added into `input_csv` directory
- First start pyspark by running
`make age`
- In the second terminal
`mv tmp/1.csv input_csv/`
- Observe average age computed in the first terminal
- Next
`mv tmp/2.csv input_csv/`
- Notice: it is important to use `mv` because files for Spark Streaming are supposed to be immutable. `mv` simply renames a file instantaneously while something like `cp` might take time. As a result Spark might process the first part of the file it notices and ignore the rest.