

Introduction

Igor Yakushin
`ivy2@uchicago.edu`

December 27, 2019

Introduction

- Apache Spark is a fast and general-purpose cluster computing system.
- It provides high-level APIs in Java, Scala, Python and R
- Typically all the functionality is available in Scala and Java while APIs in other languages might be behind. For example, Spark's GraphX library is only available in Scala.
- It also supports a rich set of higher-level tools including Spark SQL for SQL and structured data processing, MLlib for machine learning, GraphX for graph processing, and Spark Streaming.
- Spark can run within Hadoop cluster and knows how to deal with various Hadoop components (like HDFS, Hive, HBase, etc) and data formats.
- Spark can also run on a standalone computer or a regular HPC cluster like midway.
- Spark can utilize multiple nodes and multiple CPU cores in a node.

Introduction

- Python API, **pyspark**, can be used
 - in batch mode
 - interactively
 - in a python command prompt
 - in Jupyter notebook
- While **MLlib** - Spark's machine learning library - has some Neural Network routines, for more advanced Deep Learning framework consider using Spark in combination with **BigDL** library from Intel that knows how to work with Spark's RDDs

Introduction

- Spark was initially started by Matei Zaharia at UC Berkeley's AMPLab in 2009, and open sourced in 2010 under a BSD license.
- In 2013, the project was donated to the Apache Software Foundation and switched its license to Apache 2.0.
- Michael Franklin, who co-founded and directed AMPLab when Spark was created, is now professor and chair at the Computer Science Department of the University of Chicago.
- The current version of Spark is 2.4.4.
- The upcoming version 3.0.0 promises better integration with GPUs and deep learning frameworks, much faster performance

Introduction

- Before Spark 2.0, the main abstraction of Spark was the Resilient Distributed Dataset - RDD
- After Spark 2.0, RDD is replaced by DataFrame
- RDDs are still supported
- It is recommended now to use DataFrames instead of RDDs:
 - provides SQL interface to data - more convenient to program
 - much faster since a query optimization, similar the one used for SQL in RDBM, is applied to queries on DataFrames but not on RDDs
- We shall cover both since Spark is still in the state of transition. For example:
 - There are two streaming libraries:
 - Structured Streaming - based on DataFrames,
 - DStreams - based on RDDs
 - There are two APIs to machine learning library:
 - the old one based on RDD is in a maintenance state - only bug fixes are applied to it but no new features are introduced;
 - the new machine learning library based on DataFrames is still catching up with the functionality of the old library